Department of Computer Science Technical Reports

Department of Computer Science

1990

Convex Decomposition of Polyhedra and Robustness

Chanderjit Bajaj

Tamal K. Dey

Report Number:

90-990

CONVEX DECOMPOSITION OF POLYHEDRA AND ROBUSTNESS*

Chanderjit Bajaj
and
Tamal K. Dey

Computer Sciences Department
Purdue University
Technical Report CSD-TR-990
CAPO Report CER-90-27
June, 1990

Convex Decomposition of Polyhedra and Robustness[*]

Chanderjit L. Bajaj   *Tamal K. Dey*

Department of Computer Science
Purdue University
West Lafayette, IN 47907

Abstract

We present a simple algorithm to compute a convex decomposition of a non-convex, non-manifold polyhedron of arbitrary genus (handles). The algorithm takes a non-convex polyhedron with n edges and r *notches* (features causing non-convexity in the polyhedra) and produces a worst-case optimal $O(r^2)$ number of convex polyhedra S_i, with $\bigcup_i S_i = S$, in $O(nr^2)$ time and $O(nr)$ space. Recently, Chazelle and Palios have given a fast $O(n\ r\ +\ r^2 log r)$ time algorithm to tetrahedralize a non-convex simple polyhedron. Their algorithm, however, works for a simple polyhedron of genus 0 and with no shells (inner boundaries). The input polyhedron of our algorithm may have arbitrary genus and inner boundaries and may be a non-manifold. We also present an algorithm for the same problem while doing only finite precision arithmetic computations.

1   Introduction

The main purpose behind decomposition operations is to simplify a problem for complex objects into a number of subproblems dealing with simple objects. In most cases a decomposition, in terms of a finite union of disjoint convex pieces is useful and this is always possible for polyhedral models [4, 8]. Convex decompositions lead to efficient algorithms, for example, in geometric point location and intersection detection, see [8]. Our motivation stems from the use of geometric models in SHILP, a solid model creation, editing and display system being developed at Purdue [2]. Specifically, a disjoint convex decomposition of simple polyhedra allows for more efficient algorithms in motion planning, in the computation of volumetric properties, and in the finite element solution of partial differential equations. In what follows, we use the following definitions. The surface of a polyhedron S is called a 2-manifold if for each point on the surface of S, there exists an $\epsilon - neighborhood$ which is homeomorphic to a 1-sphere or a circle [19]. Polyhedra, which have 2-manifold surface are called manifold polyhedra. Polyhedra which are not manifold are called non-manifold polyhedra. Non-manifold polyhedra may have incidences as illustrated in the Figure 1. Manifold polyhedra with holes are homeomorphic to toruses with one or more handles. Manifold polyhedra with inner boundaries are homeomorphic to 3-dimensional annuli i.e., spheres with bubbles inside them. A *reflex edge* of a polyhedron is the one where the inner dihedral angle subtended by two incident facets is greater than 180°.

Related Work: The problem of partitioning a non-convex polyhedron S into a *minimum* number of convex parts is known to be NP-hard [16, 18]. Rupert and Seidel [20] also show that the problem of determining whether a non-convex polyhedron can be partitioned into tetrahedra, without introducing Steiner
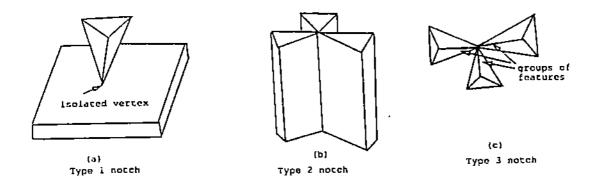

1

Figure 1: Non-manifold incidences or *special notches*.

points, is NP-hard. For a given polyhedron S with n edges of which r edges are *reflex*, Chazelle [4, 5] established a worst-case, $O(r^2)$ time lower bound on the complexity of the decomposition problem, allowing Steiner points, and gave an algorithm that produces a worst-case, optimal number $O(r^2)$ convex polyhedra in $O(nr^3)$ time and $O(nr^2)$ space. Recently, Chazelle and Palios [6], also gave an $O(nr + r^2 log r)$ time algorithm to tetrahedralize a subclass of non-convex polyhedra. The allowed polyhedra are all homeomorphic to a 2-sphere, i.e., have no holes(genus 0) and shells (inner boundaries) and are manifold.

Results: In section 3, we first present an algorithm to compute a disjoint convex decomposition of a manifold polyhedron S which may have an arbitrary number of holes and shells. Given such a polyhedron S with n edges of which r are *reflex*, the algorithm produces a worst case optimal $O(r^2)$ number of convex polyhedra S_i, with $\bigcup_i S_i = S$ in $O(nr^2)$ time and $O(nr)$ space. We extend this algorithm to non-manifold polyhedra which may not have abutting edges or facets but may have incidences as illustrated in Figure 1. In section 4, we give an algorithm for the same problem, which uses sophisticated heuristics based on geometric reasoning to overcome the inaccuracies involved with finite precision arithmetic computations. This algorithm runs in $O(nr^2 + nr log n + r^3 log n + r^4)$ time and in $O(nr)$ space.

2 Preliminaries

2.1 Data Structure and Definitions

Let S be a polyhedron, possibly with holes and shells, and having s vertices : $\{v_1, v_2, ..., v_s\}$, n edges : $\{e_1, e_2, ..., e_n\}$ and q facets : $\{f_1, f_2, ..., f_q\}$.

Polyhedron Data Structure: The polyhedron S with arbitrary number of holes and shells, is represented by a collection of vertices, edges, and facets, each of which is maintained as structures similar to the representations of [15].
Vertices: Each vertex is represented with two fields.

1. *vertex.coordinates*: contains the three dimensional coordinates of the vertex.

2

2. *vertex.adjacencies*: contains pointers to the edges incident on the vertex.

Edges: Each edge is represented with two fields.

1. *edge.vertices*: contains pointers to the incident vertices.

2. *edge.orientededges*: contains pointers to the structures called *orientededges* which represent different orientations of an edge on each face incident on it. The orientation of an edge on a facet f is such that a traversal of the oriented edge has facet f to its right.

Orientededges: Each Orientededge is represented with three fields.

1. *orientededge.edge*: Contains pointer to the corresponding edge.

2. *orientededge.facet*: Contains pointer to the facet on which the *orientededge* is incident.

3. *orientededge.orientation*: Contains information about the orientation of the edge on the facet.

Facets: Each facet is represented with two fields.

1. *facet.equation*: contains the equation of the plane supporting the facet.

2. *facet.cycles*: contains pointers to a collection of oriented edge cycles bounding the facet. The traversal of each oriented edge cycle always has the facet to the right. Each edge cycle is represented as a linked list of structures representing the *orientededges* on the cycle. If there is a vertex touching the face, (Figure 1(a)) called an *isolated vertex*, a pointer to the vertex is included in *face.cycles* as a degenerate edge cycle.

The intersection of S with a plane P is, in general, a set of simple polygons, possibly with holes. If G is a simple polygon with vertices $v_1, v_2, ..., v_k$ in clockwise order, a vertex v_i is a *reflex vertex* of G if the inner angle between the edge (v_{i-1}, v_i) and (v_i, v_{i+1}) is $> 180°$. The vertices which are not *reflex vertices* are called *normal vertices* of G. The boundary of a polygon G can be partitioned into x-monotone maximal pieces called *monotone chains*, i.e., vertices of a *monotone chain* have x-coordinates in either strictly increasing or decreasing order. See Figure 2.

In general, non-manifold polyhedra have nonconvexity due to the following four types of features called *notches*.

1. *Type 1 notches*: These *notches* are caused by vertices which touch a face as illustrated in the Figure 1(a). The vertex on the face is called an *isolated vertex*.

2. *Type 2 notches*: More than two facets may be incident on an edge e_i as illustrated in the Figure 1(b). Two adjacent facets around the edge e_i which do not enclose any volume of S causes the nonconvexity or a *notch*. If there are $2k$ $(k > 1)$ facets incident on e_i, they form k *notches*.

3. *Type 3 notches*: These *notches* are caused by vertices where two or more *groups* of features (facets, edges) touch each other as illustrated in the Figure 1(c). The features within a *group* are reachable from one another while remaining only on the surface of S and not crossing the vertex. Actually, *type 1 notches* are a subclass of these *notches*. For convenience in the description, we exclude *type 1 notches* from the class of *type 3 notches*. The number of *groups* attached to the vertex determines the number of *type 3 notches* associated with that vertex.

4. *Type 4 notches*: An edge g of polyhedron S is a *type 4 notch* if the inner dihedral angle γ between two incident facets of g, is greater than $180°$. Nonconvexity in a manifold polyhedron S, is a result of the presence of these *notches* which are also called *reflex edges*.
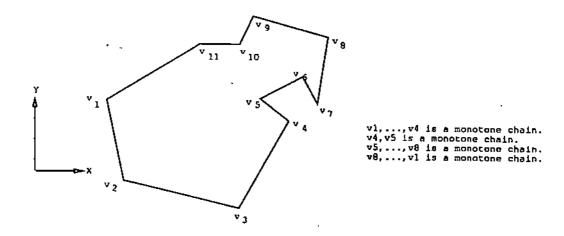
3

Figure 2: *Monotone chains* in a polygon.

The *notches* of type 1, type 2, type 3 are called *special notches* which are present only in non-manifold polyhedra. Our algorithm, first, removes all *special notches* from S creating manifold polyhedra and then proceeds in removing all *notches* of *type* 4 of the manifold polyhedra, by repeatedly cutting and splitting them with planes containing the *notches*. If an edge g is a *notch* in a manifold polyhedron, with f_g^-, f_g^+ as its incident facets, a plane P_g which contains the *notch* g and subtends an inner-angle greater than $\gamma - 180°$ with both f_g^- and f_g^+, is a valid plane which resolves the *notch* g. The chosen plane P_g is also called the *notch plane* of g. Clearly, for each *notch* g, there exist infinite choices for P_g. Note that P_g may intersect other *notches*, thereby producing *subnotches*. See Figure 3.

2.2 Useful Lemmas

In the next sections we use the following Lemmas.

As discussed in [5], one can always produce a worst case optimal number $(O(r^2))$ convex polyhedra by carefully choosing the *notch planes*.

Lemma 2.1: A manifold polyhedron S with r *notches*, can be decomposed into $\frac{r^2}{2} + \frac{r}{2} + 1$ convex pieces if all *subnotches* of a *notch* are eliminated by a single *notch plane*. Further, this convex decomposition is worst-case optimal since there exists a class of polyhedra which cannot be decomposed into fewer than $O(r^2)$ convex pieces.

Proof: See [5].

Lemma 2.2: Let G be a simple polygon with r *reflex vertices*, then the number of *monotone chains* C_s in G is bounded as $C_s \leq 6(1 + r)$.
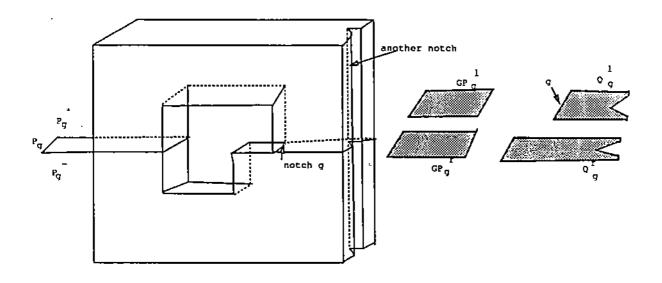**Proof:** Follows from Theorem 3, page 22 of [4]. ♣

4

Figure 3: A *notch* and its *notch plane, cross sectional map, cut.*

Below, we use the definitions from section 2.1 of *reflex, normal vertices* and *monotone chains* of a polygon.

Lemma 2.3: Let G be a simple polygon with s *normal vertices.* There are at most $O(s)$ *monotone chains* in G.

Proof: Let B be the boundary obtained by removing a vertex v and an $\epsilon - ball$ around v from the boundary of G. Add 6 more edges to B as shown in the Figure 4 to construct a new polygon G'. The polygon G' is of opposite orientation to G. Note that the vertex v always exists such that the construction of G' is possible. In fact, any vertex which is on the convex hull of the vertices of G can be taken as v. The *normal vertices* of G are the *reflex vertices* of G' except v. Moreover, constant number of edges are added to construct G' from G. Thus, G' has $O(s)$ *reflex vertices.* According to Lemma 2.2, G' has $O(s)$ *monotone chains.* The polygon G cannot have more *monotone chains* than G', which implies that G has $O(s)$ *monotone chains.*♣

In the following Lemma, the line segments of a line which are interior to a polygon are called *chords*.

Lemma 2.4: Let G be a simple polygon (possibly with holes) with r *reflex vertices.* No line can intersect G in more than $r + 1$ *chords* and $2r + 2$ points.

Proof: The proof proceeds inductively. The case for $r = 0$ is trivial. In the general step, consider a polygon G with $r = k \geq 1$ *reflex vertices.* Take an arbitrary *reflex vertex,* and resolve it by a cut through it. The cut may separate G into two polygons G_1 and G_2, of r_1 and r_2 *reflex vertices* respectively, such that
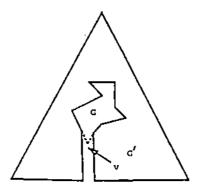
5

Figure 4: Constructing a polygon of opposite orientation.

$r_1 + r_2 \leq k - 1$. Furthermore, the number of *chords* in G cannot exceed the sum of the number of *chords* in G_1 and G_2. Therefore, using the induction hypothesis, one can conclude that L intersects G in no more than $r_1 + 1 + r_2 + 1 \leq k + 1$ *chords*. If, however, the cut does not split G, one ends up with a polygon G' of at most $k - 1$ *reflex vertices*. Since the line L may intersect the cut, just performed, the number of *chords* in G is less than or equal to that in G', which again implies that the former is less than or equal to $k - 1 + 1 \leq k + 1$. ♣

2.3 Nesting of Polygons

The following polygon nesting problem arises as a subproblem in our polyhedral decomposition.

Problem: Let $\wp$ be a set of k simple polygons $G_i, i = 1..., k$ which do not intersect along their boundaries. Corresponding to each polygon G_i we define $ancestor(G_i)$ as the set of polygons containing G_i. The polygon G_k in $ancestor(G_i)$ is called the *parent* of G_i if $ancestor(G_k) = ancestor(G_i) - G_k$. Notice that there may not exist any such G_k since $ancestor(G_i)$ may be empty. In that case, we say that the *parent* of G_i is *null*. Any polygon with *parent* G_k is called the *child* of G_k. See Figure 5. The *nesting structure* of $\wp$ is an acyclic directed graph(a forest of trees) in which there is a node n_i, corresponding to each polygon G_i in $\wp$, and a directed edge from a node n_i to n_j if and only if G_j is the parent of G_i. The polygon nesting problem is to compute the *nesting structure* of a set $\wp$ of simple nonintersecting polygons.

Lemma 2.5: The problem of polygon nesting for k simple, nonintersecting polygons can be solved in $O(s + t\log t)$ time assuming exact numerical calculations, where s is the total number of vertices and t is the total number of *monotone chains* of all input polygons.

Proof: See [3]. Though, the algorithm, given in [3], uses a slightly different type of *monotone chains*, called *subchains*, it also works for the *monotone chains* as defined in this paper. With this slight modification, Theorem 2.1 of [3] can be restated as Lemma 2.5 given above.♣
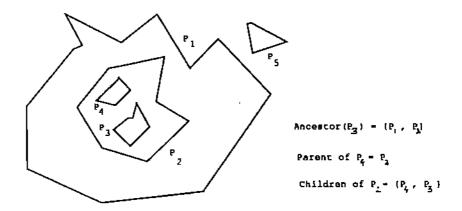
6

Figure 5: Polygon nesting.

3 Convex Decomposition

We assume the input polyhedron S to be a manifold while describing the algorithm and extend it to handle non-manifold polyhedra later. By this assumption, *notches* in S are only *reflex edges*. The algorithm for decomposing a polyhedron S with r *notches* consists of a sequence of intersections of polyhedra with *notch planes*. Hence, we first describe the method of cutting a polyhedron S by a *notch plane* P_g of a *notch* g.

3.1 Cross Sectional Map

The *notch plane* P_g: $ax + by + cz + d = 0$ defines two open half spaces P_g^+ : $ax + by + cz + d > 0$ and P_g^- : $ax + by + cz + d < 0$. The closure of P_g^+ is $P_g^u = P_g^+ \cup P_g^\ell$, where P_g^ℓ : $ax + by + cz + d = 0$ is the oriented plane P_g with normal (a, b, c) pointing into the exterior of P_g^+. Similarly, the closure of P_g^- is $P_g^b = P_g^- \cup P_g^r$ where P_g^r : $-ax - by - cz - d = 0$ is the oriented plane P_g, with normal $(-a, -b, -c)$ pointing into the exterior of P_g^-.

Cutting a polyhedron S with the plane P_g is equivalent to computing

$$S \cap P_g^u = (S \cap P_g^+) \cup GP_g^l$$
$$S \cap P_g^b = (S \cap P_g^-) \cup GP_g^r$$

where

$$GP_g^l = (closure(S \cap P_g^+)) - (S \cap P_g^+)$$
$$GP_g^r = (closure(S \cap P_g^-)) - (S \cap P_g^-).$$

We also frequently refer to GP_g^ℓ and GP_g^r as *cross sectional maps*. Note that for a polyhedron S, and a plane P_g, the *cross sectional maps* GP_g^ℓ and GP_g^r may be different. See for example, Figure 3. However, one can observe that GP_g^ℓ and GP_g^r are same if there is no facet of S lying on the *notch plane*. For simplicity, we assume GP_g^ℓ and GP_g^r to be congruent and refer to it as GP_g in describing the algorithm. With minor modifications of the algorithm one may remove this restriction.

7

The construction of GP_g corresponding to the *notch plane* P_g is the crucial part in splitting a polyhedron S to remove g. The unique polygon Q_g(possibly with holes) in GP_g, called the *cut*, supporting the *notch* g is determined and S is split along this cut. Actually, splitting S along the cut instead of the *cross sectional map*, is sufficient to remove the *notch* g of S. Note that because of this, S may not get separated into two different pieces after the split. In Figure 3, the removal of the *notch* g through the *cut* Q_g does not separate S. The *notch* g may lie on the inner or the outer boundary of Q_g. We denote the boundary containing g as B_g.

- **Step I:** Determine Q_g. This calls for computing inner and outer boundaries of Q_g.

- **Step II:** Split S. While describing the algorithm we assume S is separated into two pieces by the cut Q_g. The case where S is merely spliced by Q_g instead of getting separated into two pieces does not incur any extra overhead to our algorithm.

In what follows, we use (lower case) letters, s,u, for counting vertices, m,n,p,t for counting edges, and q for counting facets. Let S have p edges of which r are *reflex*.

3.2 Description of the Algorithm

Step I : First compute all boundaries B present in the *cross sectional map* GP_g. Visit all the facets of S in turn. If a facet f_i intersect the *notch plane* P_g, all intersection points are computed. Let $a_1^i, a_2^i, ..., a_k^i$ be the intersection points on the edges $e_1^i, e_2^i, ..., e_k^i$ respectively of f_i. These intersection points can be sorted along the line of intersection $P_g \cap f_i$ at a cost, linear in number of edges present in the facet f_i using the algorithm of [13]. Associate this sorted sequence of intersection points with f_i. Further, with each intersection point a_j^i, keep the information of the edge e_j^i. Pick an intersection point a_l^j. Continue to
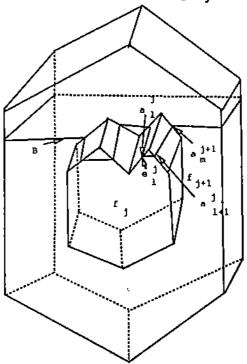

Figure 6: Computation of a boundary in the *cross sectional map*.

construct the boundary B containing a_l^j as follows. See Figure 6. One of the segments $a_l^j a_{l+1}^j$ and $a_l^j a_{l-1}^j$ lies inside f_j. Without loss of generality, assume $a_l^j a_{l+1}^j$ lies inside f_j. Join a_l^j and a_{l+1}^j and continue from a_{l+1}^j

to determine other edges of the boundary B. Consider the other facet f_{j+1} adjacent to the edge e_{l+1}^j. The facet f_{j+1} can be retrieved in constant time in our data structure. Determine the intersection point a_m^{j+1} adjacent to a_{l+1}^{j+1} on f_{j+1} such that the segment $a_{l+1}^{j+1} a_m^{j+1}$ lies inside f_{j+1}. Without loss of generality, assume a_m^{j+1} is ordered after a_{l+1}^j in the sorted sequence of intersection points associated with f_{j+1}. Note that the points a_{m-1}^{j+1} and a_{l+1}^j are same. Proceed from a_m^{j+1} and continue the above procedure until the initial point a_l^j is reached. This completes the computation of B. Once a boundary computation is completed, pick up another intersection point which has not been visited yet and construct the corresponding boundary by the above method. Continuing this procedure until all intersection points are visited gives all the boundaries present in the *cross sectional map*. The adjacent points of an intersection point on a facet can be retrieved in constant time if the sorted sequence of intersection points is maintained as a doubly linked list.

Next, determine the inner and outer boundaries of Q_g. It is trivial to determine the boundary B_g containing the notch g. One can determine whether B_g is an inner or outer boundary of Q_g by checking the orientation of the edges on the boundary. Orientation of each such edge is determined in constant time since the orientations of the *notch plane* and the facets intersecting the *notch plane* are known.

Case (i): B_g is an outer boundary of Q_g. Let I_i be any inner boundary of Q_g. The boundary I_i itself constitutes a simple polygon. Polygon I_i will have at least one (actually at least three) vertex, which is a *normal vertex*. Since I_i is the inner boundary of Q_g, the vertices which are *normal vertices* of polygon I_i are *reflex vertices* of Q_g. Definitely, *reflex vertices* of Q_g lie on *notches* of S. This implies that all inner boundaries of Q_g will have a point which is the intersection point of P_g with a *notch* of S. Determine the set W of boundaries having at least one point where a *notch* of S and P_g intersect. This takes $O(u')$ time, where u' is the number of vertices present on the cross sectional map. Call the boundaries in the set $W \cup B_g$ as *interesting boundaries*. Certainly, the number of *interesting boundaries* is $O(t)$ where t is the number of *notches* intersected by the *notch plane* P_g. The *interesting boundaries* which are outer boundaries of some polygon in the *cross sectional map*, have $O(t)$ *reflex vertices*. On the other hand, the *interesting boundaries* which are inner boundaries of some polygon in the cross sectional map have $O(t)$ *normal vertices*. Thus, according to the Lemma 2.2 and 2.3, there are at most $O(t)$ *monotone chains* in the *interesting boundaries*. If there are u vertices on the *interesting boundaries*, the inner boundaries of Q_g can be determined in $O(u + t \log t)$ time using Lemma 2.5. The computation of the boundaries in the *cross sectional map* takes at most $O(p)$ time. This is due to the fact that the sorted sequence of intersection points on each facet can be computed at a cost linear in number of edges of the facet. Thus, in this case, the inner and outer boundaries of Q_g can be determined in

$$O(p + u + u' + t \log t) = O(p + t \log t)$$

time, since $u = O(u') = O(p)$.

Case (ii): B_g is an inner boundary of Q_g. Visit all edges of the polyhedron S, being split to compute the sorted sequence of intersection points on each facet and compute all the boundaries present in the *cross sectional map*. Determine the boundaries which contain the boundary B_g inside them. Call these boundaries, together with B_g, as *interesting boundaries*. This takes $O(p + u') = O(p)$ time. Apply the polygon nesting algorithm of [3] on these *interesting boundaries* to detect the parent polygon of B_g which is the outer boundary of Q_g. The *interesting boundaries* can be partitioned into two classes according to whether they are inner or outer boundaries of some polygon. It is not hard to see that there must be as many inner boundaries as outer boundaries. Hence, the number of *interesting boundaries* is bounded above by twice the number of inner boundaries present in the *cross sectional map*. As discussed in the previous case, this number must be bounded above by the number of *notches* intersected by the *notch plane*. Thus, there are $O(t)$ *interesting boundaries*. Further, the number of *monotone chains* present in these *interesting boundaries* can be at most $O(t)$. Hence, as in the previous case, the inner and outer boundaries of Q_g can be determined in $O(p + t \log t)$ time.

Step II: Separation of S corresponding to the cut Q_g is carried out by splitting facets which are intersected by Q_g. Suppose f_i is such a facet which is to be split at $a_1^i, a_2^i, ..., a_k^i$ which are on the edges $e_1^i, e_2^i, ..., e_k^i$. The splitting of f_i consists of splitting the edges on which $(a_1^i, a_2^i, ..., a_k^i)$ lies. Visit only the intersection points corresponding to the vertices of Q_g and for each such intersection point spend constant time for setting relevant pointers to carry out the split operation. Create two oppositely oriented facets at the same geometric location corresponding to the cut Q_g. Adjust all the modified incidences properly. A depth first traversal in the modified vertex list completes the separation of S by collecting all the pertinent features of each piece. This process cannot take more than $O(p)$ time. Combining the costs of *StepI* and *StepII* yields the following Lemma.

Lemma 3.1. A manifold polyhedron S of genus 0, having p edges can be partitioned with a *notch plane* P_g of a *notch* g in $O(p + t\log t)$ time and in $O(p)$ space where t is the number of *notches* intersected by P_g.
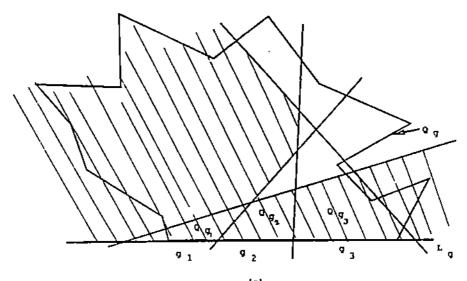
We can generalize the above result for a polyhedron of arbitrary genus. For this, as described in [5], we need to handle the situation when the cut does not separate S into two pieces, but only creates two new facets supporting the cut at the same geometric location. A depth-first search in the vertex list determines whether the cut separates S into two pieces or not.

Lemma 3.2. Let $S_1, S_2, ..., S_k$ be the polyhedra in the current decomposition, where each S_i contains a *subnotch* g_i of a *notch* g of a manifold polyhedron S with n edges and r *notches*. Let m_i and u_i be the number of edges and vertices on Q_{g_i} respectively. Then m and u, the total number of edges and vertices on all the cuts supported by the *subnotches* of the *notch* g are given as $m = \sum_{i=1}^{k} m_i = O(n)$ and $u = \sum_{i=1}^{k} u_i = O(n)$.

Proof: Consider the cut Q_g produced by the intersection of S with P_g. The region in Q_g is divided into smaller facets by *notch lines* produced by the intersection of other *notch planes* with P_g. We focus on the facets $Q_{g_1}, Q_{g_2}, ..., Q_{g_k}$ adjacent to the *subnotches* $g_1, g_2, ..., g_k$ of the *notch* g.

Consider the set of *notch lines* which divides Q_g and the line L_g corresponding to the *notch* g. They produce a line arrangement [8] on the *notch plane* P_g. Consider the facets adjacent to the line L_g in this arrangement. These are called *zones* of L_g. See Figure 7(a). Let us denote the set of these *zones* by Z_g and their vertices and edges by V_g and E_g respectively. It is proved that (Theorem 5.3,pp. 89, [8]) $|V_g| \leq 5l - 3$ and $|E_g| \leq 5l - 1$ if there are l lines in the arrangement. Overlaying Q_g on Z_g produces $Q_{g_1}, Q_{g_2}, ..., Q_{g_k}$. Let V_g' and E_g' denote the set of vertices and edges in $Q_{g_1}, Q_{g_2}, ..., Q_{g_k}$. The vertices in V_g' can be partitioned into three different sets, namely, T_1, T_2, T_3. The set T_1 consists of vertices formed by the intersections of *notch lines*, T_2 consists of vertices formed by the intersections of edges of Q_g and T_3 consists of vertices formed by the intersections of *notch lines* and edges of Q_g. Certainly, $|T_1| \leq |V_g| = 5l - 3$ since overlaying of Q_g on Z_g cannot introduce any vertices in T_1. If Q_g has u' edges, $|T_2| \leq u'$.

To count the number of vertices in T_3, consider an edge e in E_g which contributes one or more segments e_s to E_g' as a result of intersections with Q_g. There must be at least one *reflex vertex* of Q_g, present between two successive edge segments e_s. Charge a cost of 1 to the *reflex vertex* which lies to the left (or, right) of each segment and charge a cost of 1 to e itself for the leftmost (or, rightmost) segment. We claim that each *reflex vertex* of Q_g is charged at most once by this method. Suppose, on the contrary, a *reflex vertex* is charged twice by this procedure. Then, that *reflex vertex* must appear between two segments of two edges in E_g as shown in Figure 7(b). As can be easily observed, all four edge segments cannot be adjacent to the regions incident on an edge g of Q_g. This contradicts our assumption that all these four segments are present in E_g'. Hence, the total charge incurred upon the *reflex vertices* of Q_g and the edges of E_g can be at most $r_g + 5l - 1$ where r_g is the number of *reflex vertices* present in Q_g.
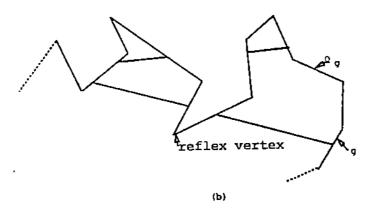
10

(a)

Shaded regions are zones of L_g.

(b)

Figure 7: *Zones* of a line and *cuts*.

This implies that as a result of intersections with Q_g, at most $r_g + 5l - 1$ segments of edges in E_g are contributed to E'_g. Hence, $|T_3| \leq 2(r_g + 5l - 1)$. Putting all these together, we have

$$
\begin{aligned}
|V'_g| &= |T_1| + |T_2| + |T_3| \\
&\leq 5l - 3 + u' + 2r_g + 10l - 6 \\
&\leq 15l + u' + 2r_g - 9.
\end{aligned}
$$

Since there can be at most r *notch planes*, $l \leq r$. Certainly, $r_g \leq r$ and $u' \leq n$. This gives

$$
u = |V'_g| \leq 15r + n + 2r - 9 = O(n + r) = O(n).
$$

Since $Q_{g_1}, Q_{g_2}, ..., Q_{g_k}$ form a plane graph, we have

$$
m = |E'_g| = O(|V'_g|) = O(n).\clubsuit
$$

Lemma 3.3: The total number of edges in the final decomposition of the polyhedron S with r *notches* and n edges is $O(nr)$.

Proof: Total number of edges in the final decomposition consists of newly generated edges by the cuts, and the edges of S which are not intersected by any *notch plane*. Since the total number of edges present in all the cuts corresponding to a *notch* is $O(n)$, the total number of newly generated edges by each *notch plane* is $O(n)$. Thus r *notch planes* generate $O(nr)$ new edges. Hence, the total number of edges in the final decomposition is $O(nr + n) = O(nr).\clubsuit$

Theorem 3.1: A manifold polyhedron S, possibly with holes and shells and having r *notches* and n edges can be decomposed into $O(r^2)$ convex polyhedra in $O(nr^2)$ time and $O(nr)$ space.

Proof: Decomposition of a polyhedron consists of a sequence of cuts through the *notches* of S. Assign a *notch plane* for each *notch* in S in $O(r)$ preprocessing time. Remove each *notch* by removing all of its *subnotches* with the *notch plane* assigned to this *notch*. Each planar cut to remove a *subnotch* in a polyhedron, can be carried out by the method described above. According to the Lemma 2.1 this produces $O(r^2)$ convex pieces at the end since all *subnotches* of a *notch* are removed by a single *notch plane*.

Worst-Case Complexity: At a generic instance of the algorithm, let $S_1, S_2, ..., S_k$ be the k distinct (non-convex) polyhedra in the current decomposition, where each S_i contains the *subnotch* g_i of a *notch* g which is going to be removed. Let S_i have p_i edges and $p = \sum_{i=1}^{k} p_i$. Let t_i be the number of *notches* intersected by P_{g_i} in S_i and $t = \sum_{i=1}^{k} t_i$.

Applying Lemma 3.1, each *subnotch* g_i in S_i can be removed in $O(p_i + t_i log t_i)$ time and in $O(p_i)$ space. Thus, removal of a *notch* g can be carried out in $O(\sum_{i=1}^{k}(p_i + t_i log t_i))$ time and in $O(\sum_{i=1}^{k} p_i)$ space. By Lemma 3.3, $\sum_{i=1}^{k} p_i = O(nr)$. Since a *notch plane* can intersect at most $r - 1$ *notches*, $t = O(r)$. This gives,

$$
\sum_{i=1}^{k} t_i\, log t_i = O(t\, log t) = O(r log r).
$$

Hence, a *notch* g can be removed in $O(nr + r log r) = O(nr)$ time. Thus, elimination of r *notches* takes $O(nr^2)$ time. In Lemma 3.3, we prove that the total number of edges in the final decomposition of S is $O(nr)$. This implies that the space complexity of polyhedral decomposition is $O(nr)$. $\clubsuit$

3.3   Decomposition of non-manifold polyhedra

For a non-manifold polyhedron S, nonconvexity results from *notches* of four types as discussed in section 2.1. Let S have n edges and r *notches*. Preprocess S as follows to remove the *notches* of first three types, called the *special notches*. Let this process produce a decomposition $S_1, S_2, ..., S_l$ where each S_i is a manifold polyhedron having notches of only the fourth type. Apply Theorem 3.1 on each of them to obtain a worst-case optimal convex decomposition.

Removal of type 1 notches: In this case, as can be observed from the Figure 1(a), the vertex v_i causing the nonconvexity is detached from the facet f_i on which it is incident as *isolated vertex*. Identifying these vertices and detaching them from corresponding facets take at most $O(n)$ time.

Removal of type 2 notches: In this case, more than two facets are incident on an edge e_i. Let these facets be $f_1, f_2, ..., f_{r_i}$. Consider a cross section C which is the intersection of the facets incident on e_i with the plane P perpendicular to the edge e_i. C consists of edges $e_j = (f_i \cap P)$. Sort the facets circularly around the edge e_i by the circular sort of the edges e_j which are incident on $e_i \cap P$. Pair the adjacent facets which enclose a volume of S. Let this pairing be $(f_1, f_2), (f_3, f_4), ..., (f_{r_i-1}, f_{r_i})$. Create an edge between each pair of facets and delete the edge e_i. All these edges are at the same geometric location of e_i. Adjust all the incidences properly. Sorting of facets around the edge e_i takes $O(r_i log r_i)$ time. The adjustment time for of all incidences in the internal representation of S cannot exceed $O(n)$. Thus, removal of all *type 2 notches* takes at most $O(r log r + nr) = O(nr)$ time.

Removal of type 3 notches: Let v_i be a vertex which corresponds to a *type 3 notch*. In this case, collect the features (edges,facets) incident on v_i which are reachable from one another while remaining always on the surface of S and never crossing v_i. This gives a partition of the features incident on v_i into smaller groups. For each such group, create a vertex at the same geometric location of v_i and adjust all the incidences properly. This in effect, removes the nonconvexity caused by v_i. All such vertices causing *type 3 notches* in S can be identified in $O(n)$ time by depth first traversal in the underlying graph of S. Removal of each such notch takes at most $O(n)$ time. Thus, all *type 3 notches* can be removed in $O(nr)$ time.

Removal of all the above three types of *notches* generates at most $O(n)$ new edges and produces at most k manifold polyhedra where k is the number of *special notches* in S.

Theorem 3.2: A non-manifold polyhedron S, possibly with holes and shells and having r *notches* and n edges can be decomposed into $O(r^2)$ convex polyhedra in $O(nr^2)$ time and $O(nr)$ space.

Proof: Remove all *special notches* from S in $O(nr)$ and $O(n)$ space as discussed above. Let $S_1, S_2, ..., S_l$ be the manifold polyhedra created by this process. Let S_i have n_i edges of which r_i are reflex. Using Theorem 3.1 on each of them, we conclude that S can be decomposed into $O(r^2)$ convex polyhedra in $O(\sum_{i=1}^{l} n_i r_i^2) = O(nr^2)$ time and in $O(\sum_{i=1}^{l} n_i r_i) = O(nr)$ space.

Decomposition into Tetrahedra: Let $S_1, S_2, ..., S_k$ be the convex polyhedra produced by convex decomposition of a polyhedron S. Each convex piece with p_i edges can be triangulated into $O(p_i)$ tetrahedra in a straightforward manner (triangulate every convex facet and then tetrahedralize by choosing a point in the interior of the convex polyhedra). This takes at most $O(p_i)$ time for each convex piece. Hence, triangulation of all pieces takes $O(\sum_{i=1}^{k} p_i) = O(nr)$ time producing $O(nr)$ tetrahedra.

13

4 Convex Decomposition under Finite Precision Arithmetic

Motivation: When implementing geometric operations stemming from practical applications, one cannot ignore the degenerate geometric configurations that often arise, as well as the need to make specific topological decisions based on imprecise finite precision numerical computations [12, 15, 23]. We model the inexact arithmetic computations by ε-arithmetic [10, 11] where the arithmetic operations $+, -, \div, \times$ are performed with relative error of at most ε. Under this model, the absolute error in distance computations of one polyhedral feature from another is bounded by a certain quantity $\delta = k\varepsilon B$, where B is the maximum value of any coordinate and k is a constant. See [17]. When making decisions about the incidence of these polyhedral features (vertices, edges, facets), on the the basis of the computed distance(with sign), one can rely on the sign of the computation only if the distance is greater than δ. On the other hand, if the computed distances are less than δ, one also need to consider the topological constraints of the geometric configuration to decide on a reliable choice. In particular, in regions of uncertainity i.e. within the δ-ball, the choices are all equally likely that the computed quantity, is negative, zero or positive. Such decision points of uncertainity, where several choices exist, are either "independent" or "dependent". At independent decision points, any choice may be made from the finite set of local topological possibilities, while the choice at dependent decision points should ensure that it does not contradict any previous topological decisions. The algorithm which follows this paradigm would never fail, though it may not always compute a valid output. Such algorithms have been termed as parsimonious by Fortune [10].

An algorithm under ε-arithmetic, is called robust if it computes an output which is exact for some perturbed input. It is called stable if the perturbation required is small. Recently, in [10, 11, 17] authors have given robust and stable algorithms for some important problems in two dimensions. Except [14], there is no known robust algorithm for any problem in three dimensions. The difficulty arises due to the fact that the perturbations in the positions of the polyhedral features may not render a valid polyhedron embedded in $\Re^3$. In [14], Hopcroft and Kahn discuss the existence of a valid polyhedron which admits the positions of the perturbed vertices of a convex polyhedron. The case of non-convex polyhedra is perceived to be hard and requires understanding the deep interactions between topology and perturbations of polyhedral features of non-convex polyhedra.

Karasick [15] gives an algorithm for the problem of polyhedral intersection where he uses geometric reasoning to avoid conflicting decisions about polyhedral features. In this paper, we extend the results in [15] and provide an algorithm for the problem of polyhedral decomposition which also uses geometric reasoning to avoid conflicting decisions. Though, as yet we are unable to prove our algorithm to be parsimonious, we report various heuristics we have implemented in our effort to make the decomposition algorithm robust and stable. We also describe a worst case running time bound for the algorithm under the ε-arithmetic model.

More related work: The issue of robustness in geometric algorithms have recently taken added importance because of the increasing use of geometric manipulations in computer-aided design, and solid modeling [1]. Edelsbrunner and Mucke [9], and Yap [24], suggest using expensive symbolic perturbation techniques for handling geometric degeneracies. Sugihara and Iri [23], and Dobkin and Silver [7], describe an approach to achieve consistent computations in solid modeling, by ensuring that computations are carried out with sufficiently higher precision than used for representing the numerical data. There are drawbacks however, as high precision routines are needed for all primitive numerical computations, making algorithms highly machine dependent. Furthermore, the required precision for calculations is difficult to a priori estimate for complex problems. Segal and Sequin [21] estimate various numerical tolerances, tuned to each computation, to maintain consistency. Milenkovic [17] presents techniques for computing the arrangements of a set of lines in two dimensions robustly. He introduces the concept of *pseudo lines* which preserves some basic topological properties of lines and computes the arrangements in terms of these *pseudo lines*. Hoffmann,

Hopcroft and Karasick [12], and Karasick [15], propose using geometric reasoning and apply it to the problem of polyhedral intersections. Sugihara [22] uses geometric reasoning to avoid redundant decisions and thereby eliminate topological inconsistencies in the construction of planar Voronoi diagrams. Guibas, Salesin and Stolfi [11] propose a framework of computations called ε-geometry, in which they compute an exact solution for a perturbed version of the input. So does Fortune [10] who applies it to the problem of triangulating two dimensional point sets.

4.1 Intersection & Incidence Tests

In what follows, we assume the input polyhedra are manifold. Non-manifold polyhedra can be handled as discussed in the previous section. It is clear from discussion of our previous algorithm that numerical computations are needed in different types of intersections and incidence testings. We assume minimum feature criteria for polyhedra as follows. The distance between two distinct vertices or between a vertex and an edge and the dihedral angle between any two facets may not be less than a minimum value. The choice of this minimum threshold value is described in our algorithm. To decide whether an edge is intersected by a plane, one must decide the classifications of its terminal vertices with respect to the same plane. The same classification of a vertex is used to decide the classification of all the features incident on that vertex. This, in effect, avoids conflicting decisions about the polyhedral features. The decisions about different types of intersections and incident testings are carried out by three basic tools, namely, (i) vertex-plane classifications, (ii) facet-plane classifications and (iii) edge-plane classifications.

The order of classifications is (i) followed by (ii) followed by (iii). Edge-plane classifications are done only after vertex-plane and facet-plane classifications. In what follows, assume the equation of a plane $P_i : a_i x + b_i y + c_i z + d$ is normalized with $a_i^2 + b_i^2 + c_i^2 = 1$.

Vertex-Plane Classification: To classify the incidence of a vertex $v_i = (x_i, y_i, z_i)$ w.r.t the plane $P : ax + by + cz + d = 0$, compute the normalized algebraic distance of v_i from P by computing $ax_i + by_i + cz_i + d$. The *sign* of this computation, viz., zero, negative, or positive, classifies v_i as "on" P (zero), "below" P (negative) or "above" P (positive), where "above" is the half space containing the plane normal (a, b, c). Accept the sign of the computations as correct if the above distance of v_i from P is larger than δ. Otherwise, apply geometric reasoning rules, as detailed below, to classify vertex v_i w.r.t. the plane P. In the following algorithmic version of the vertex-plane classification, the intersection between an edge e incident on v_i and the plane P is computed as follows. Let e be incident on planes P_1, P_2, where $P_i : a_i x + b_i y + c_i z + d_i = 0$. Compute the intersection point r of e and the plane P by computing the solution of the linear system,

$Ar = d$ where $A = \begin{bmatrix} a & b & c \\ a_1 & b_1 & c_1 \\ a_2 & b_2 & c_2 \end{bmatrix}$ $d = [-d, -d_1, -d_2,]^T$. The linear system is solved using Gaussian

elimination with scaled partial pivoting and iterative refinement.

Vertex-Plane-Classif (v_i, P)
begin
    Let $v_i = (x_i, y_i, z_i)$ be a vertex incident on edges $e_1 = (v_i, w_1), e_2 = (v_i, w_2), ..., e_k = (v_i, w_k)$.
    Let $P : ax + by + cz + d = 0$.
    Compute $l = ax_i + by_i + cz_i + d$.
    If $|l| > \delta$ then   (*Comment: Unambiguously decide via the sign of distance computation*)
        if $l > 0$ then
            classify v_i as "above"
        else
            classify v_i as "below"

else

  loop

  (**Comment:* If distance computation does not yield an unambiguous
  classification for the vertex with respect to the plane, ensure that
  the "above", "below" classification is consistent with all edges
  incident on that vertex. If such consistency cannot be ensured *then*
  the vertex is classified as "maybeon" and left for future $facet - plane$
  classifications to decide its classification consistently.*)

  Search for an edge e_i incident on v_i such that $r = e_i \cap P$ is at a distance
  greater than δ from v_i and $w_i = (x_j, y_j, z_j)$.
  Get the classification of w_i if it is already computed.
  Otherwise, compute $l' = ax_j + by_j + cz_j$.
  if $|l'| > \delta$ then classify w_i accordingly.
  if the classification of w_i is "below" or "above" *then*

    if r is in between v_i and w_i *then*

      classify v_i oppositely to that of w_i

    else

      classify v_i same as that of w_i

    endif

  endif

  endloop

  if no such edge e_i is found *then*

    classify v_i as "maybeon"

    (**Comment:* To be classified later in the facet-plane classifications *)

  endif

  endif

end.

Facet-Plane Classification: If a facet f_i is intersected by a plane P in such a way that f_i does not lie on
P then the points of intersection should necessarily be (i) collinear with the line of intersection of f_i and P,
and (ii) all the vertices of f_i on one side of the intersection line, should all be of the same classification w.r.t.
the plane P. Vertices which have been temporarily classified as "maybeon", are classified in a consistent
way, i.e., they satisfy the above two properties (i) and (ii), with perturbations of at most δ. An algorithmic
version of the facet-plane classification is given below.

Facet-Plane-Classif (f_i, P)
begin

  case

    (i) All the vertices of f_i have been classified as "maybeon":
    Classify f_i as "on" the plane and change the classification of all incident vertices to "on".
    (ii) At least one vertex v_u of f_i has been classified as "above", or "below", but no
    edge of f_i has its two vertices classified with opposite signs("below" and "above"):
    if there is only one "maybeon" vertex *then*

      classify v_i as "on" and consider v_i as $f_i \cap P$

16

else

    take two "maybeon" vertices v_i, v_j and

    classify v_i and v_j as "on".

    Let L be the line joining v_i, v_j.

    Consider L as $f_i \cap P$.

endif

loop

    for each "maybeon" vertex v_k on f_i do

    if v_k is at a distance greater than δ from L *then*

        if v_k and v_u lie on opposite sides of L *then*

            classify v_k with a classification which is opposite to that of v_u.

        else

            classify v_k with a classification which is same as that of v_u.

        endif

    endif
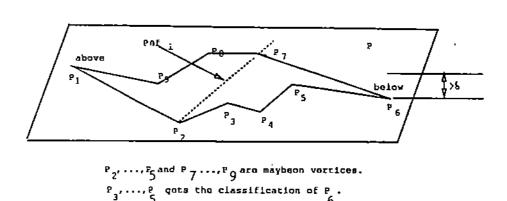
endloop

The vertices which are still not classified

$P_2, \ldots, P_5$ and $P_7 \ldots, P_9$ are maybeon vertices.

$P_3, \ldots, P_5$ gets the classification of P_6.

$P_8, \ldots, P_9$ gets the classification of P_1.

Figure 8: Case(ii) of facet-plane classification.

classify them as "on"   (*Comment:* These vertices are within a distance of δ
from L and hence will be collinear with L by a perturbation of
at most δ. See Figure 8.*)

(iii) There is an edge e whose two vertices have opposite sign classifications:

if there is no other such edge *then*

    let L be the line joining the intersection point on e and

    any "maybeon" vertex v_i.

    classify v_i as "on".

    consider L as $f_i \cap P$.

    apply methods of case (ii) to classify other "maybeon" vertices.

else

    let L be the line which fits in least square sense all the points

of intersections and apply the methods of case (ii) to classify
remaining "maybeon" vertices.
      endif
    endcase
end.

Edge-Plane Classification: An edge can get any of the three classifications which are "not-intersected", "intersected", and "on". The classifications of the vertices incident on an edge are used to classify an edge e. An algorithmic version of the edge-plane classification is given below.

Edge-Plane-Classif (e_i, P)
begin
    Let $e_i = (v_i, v_j)$.
    case
        (i) v_i and v_j are both classified as "on":
        classify e_i as "on".
        (ii) Only one of v_i, v_j, say v_i is classified as "on":
        classify e_i as "intersected" and consider v_i as $e_i \cap P$.
        (iii) v_i and v_j are classified with one as "above" and another as "below":
        classify e_i as "intersected".
        compute $r = e_i \cap P$ if it has not been computed yet.
        if r does not lie within e *then*
            choose a point at a distance of at least δ from the vertex
            which is nearest to the computed point and consider it as the intersection point of e_i and P.
        endif
        (iv) v_i and v_j are of same classifications and they are not "on":
        classify e_i as "not-intersected".
    endcase
end.

The following lemma related to consistent ordering of intersection points of a facet on the line of intersection is used in later sections.

Lemma 4.1: Let v be a vertex which is decided not to lie on the plane P and whose classification w.r.t the plane is known. Let e_1, e_2 be the edges incident on v on a facet f which are classified as "intersected". Denote the intersection points of e_1, e_2 with P as v_1 and v_2 respectively. Let O denote the ordering of v_1, v_2 on the directed intersection line $f \cap P$ which is consistent with the classification of v. If $\frac{M}{2} \geq \frac{\delta}{\sin\alpha}$ holds, O can be determined correctly. Here δ is the maximum absolute error in distance computations, α is the angle between edges e_1, e_2 on f, M is a suitably chosen large machine representable absolute value.

Proof: Consider the vertex v with incident edges e_1, e_2 on facet f. Let $L = f \cap P$ be directed as shown in Fig. 4.2 and let the actual distance of v from P be l. Suppose we know the classification of v w.r.t P. We need to determine the ordering O of v_1, v_2 on L which is consistent with the classification of v. Note

18

that the ordering of v_1, v_2 on L depends on the classification of v. See Fig 4.2(a).

Define a transformation called *max translation* as follows. Translate the plane $P : ax + by + cz + d = 0$ to $P_{maxtranslate} : ax + by + cz - M = 0$ if $d > 0$, or to $P_{maxtranslate} : ax + by + cz + M = 0$ if $d \leq 0$. Note that $P_{maxtranslate}$ is the plane P translated by the amount $M + |d|$. In the first case P is translated to its positive side and in the latter P is translated to its negative side. Let v_1', v_2' denote the intersection points of the lines containing the edges e_1, e_2 with the plane $P_{maxtranslate}$ and L' denote the directed line $P_{maxtranslate} \cap f$.

Case(i): Classification of v is same as its actual position w.r.t the plane P. See Figure 4.2(b) and 4.2(c). Transform the plane P to $P_{maxtranslate}$. If P is translated by more than l to the same side in which v lies, the ordering of v_1, v_2 is opposite to that of v_1' v_2', where l is the distance between P and v. Conversely, if P is translated by any amount to the side which does not contain v, the ordering of v_1, v_2 is same as that of v_1', v_2'.

Case(ii): Classification of v is opposite to that of its actual position w.r.t P. Transform the plane P to $P_{maxtranslate}$. If P is translated by any amount to the same side in which v has been decided to lie in, the ordering of v_1, v_2 is opposite to that of v_1', v_2'. Conversely, if P is translated by more than l to the side in which v has been decided not to lie in, the ordering of v_1, v_2 is same as that of v_1', v_2'.

In both cases, if P is translated by more than l, the ordering of v_1, v_2 can be determined from the ordering of v_1', v_2'. The ordering of v_1', v_2' can be determined exactly if the distance d' between them is greater than δ. Let l' be the distance between v and the plane $P_{maxtranslate}$. From simple geometry, one can see that $l' sin\alpha \geq \delta$ is a sufficient condition for d' to be greater than δ. P is translated by at most $l + l'$. Hence, $l + l' \leq M + |d|$. This implies

$$M + |d| \geq l + \frac{\delta}{sin\alpha}$$

is a sufficient condition for determining the ordering of v_1', v_2' exactly. Since, $min|d| = 0$ and $max|l| = \delta$, we have

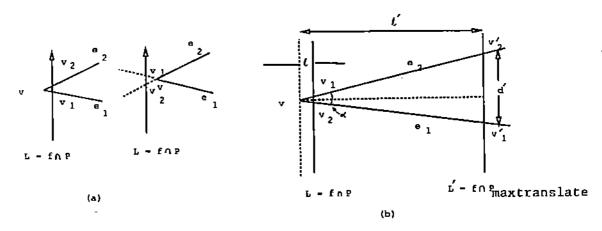$$M \geq \delta + \frac{\delta}{sin\alpha}$$

or

$$\frac{M}{2} \geq \frac{\delta}{sin\alpha}$$

is a sufficient condition for determining the ordering of v_1', v_2' exactly. The value of M is chosen to satisfy the above relation. ♣
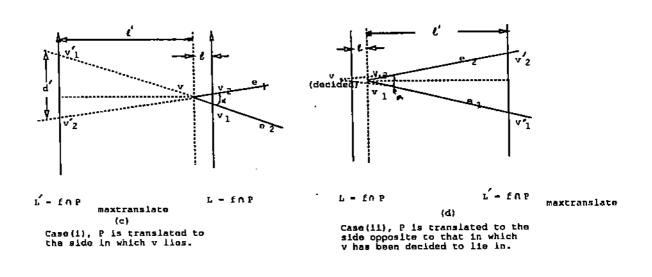

Nesting of Polygons with Finite Precision Arithmetic: The polygon nesting problem as discussed in section 2 can be solved with finite precision arithmetic if the polygons are restricted to a class of polygons called *fleshy polygons*. A polygon P is called *fleshy* if there is a point inside P such that a square with center (intersection of square's diagonals) at that point and with sides of length $64\varepsilon B$ lies inside P. B and ε have been defined earlier.

Lemma 4.2: The problem of polygon nesting for k *fleshy* polygons with s vertices and t *monotone chains* can be solved in $O(k^2 + s(t + logs))$ time under finite precision arithmetic.

Proof: See [3]. Since any vertical line (orthogonal to the x direction) can intersect at most t edges of a set of polygons having t *monotone chains*, the above time bound is obvious from the time analysis of the algorithm under finite precision arithmetic as given in [3].♣

(a)

(b)

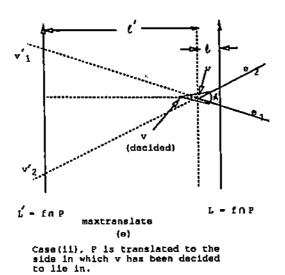Case (i), P is translated to the side opposite to that in which v lies.

(c)

Case(i), P is translated to the side in which v lies.

(d)

Case(ii), P is translated to the side opposite to that in which v has been decided to lie in.

(e)

Case(ii), P is translated to the side in which v has been decided to lie in.

Figure 9: *Maxtranslation* and Lemma 4.1.

4.2   Description of the Algorithm

The same paradigm of cutting and splitting the polyhedron about the *cuts* is followed to produce the convex decomposition of a manifold, non-convex polyhedron. Choose one of the two planes incident on a *notch* as *notch plane*. This ensures that no new planes other than facet-planes are introduced by the algorithm and thus no additional error is introduced in the plane equations containing the facets. This also guarantees that any input assumption about the planes containing the facets remain valid throughout the iterative process of cutting and splitting the polyhedron. We apply heuristics at each numerical computation through geometric reasoning to make our algorithm as parsimonious as possible. For any *notch plane* P_g the two *cross sectional maps* GP_g^l, GP_g^r are constructed and the corresponding cuts Q_g^l, Q_g^r are computed in *Step I* as detailed below. In *Step II* we split the polyhedron about these cuts which completes the removal of *notch g*.

Step I :

Constructing GP_g^l **and** GP_g^r: The edges of GP_g^l and GP_g^r are either the edges transferred from polyhedron S called *old edges*, or edges newly generated from $S \cap P_g$ called *new edges*. Note, all *new edges* will be present in both *cross sectional maps* while only some of the *old edges* may be present in either GP_g^l or in GP_g^r. As with the edges, some of the vertices of the *cross sectional maps* will be *old vertices* while some of them will be *new vertices*. To generate *old* and *new edges* on these *cross sectional maps*, compute the intersection points of each facet f with the *notch plane* using the vertex-plane, edge-plane, facet-plane classification as described before. After computing all intersection vertices (*new* and *old*) lying on the facet f, sort these vertices along the line of intersection $f \cap P_g$.
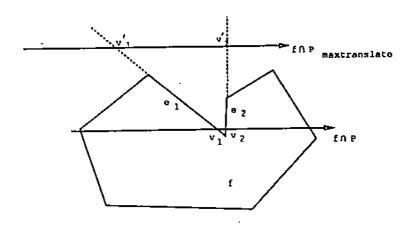

Figure 10: Consistent sorting of intersection points.

Sorting of intersection points along line $f \cap P_g$: Consider the facet f as shown in Figure 10. Let edges e_1 and e_2, incident on v intersect the plane P_g at points v_1 and v_2, both necessarily lying on line $L = f \cap P_g$. Further let v_1 and v_2 be *new vertices*. If v_1 and v_2 happen to be very close together, it may not be possible to determine their local ordering on L reliably. However, the classification of v w.r.t P_g can be used to decide this ordering consistently. Translate the plane P_g to $P_{maxtranslate}$ and compute the points

$e_1 \cap P_{maxtranslate}$ and $e_2 \cap P_{maxtranslate}$. Let these intersection points be v_1' and v_2' respectively. As the angle between edges e_1 and e_2 cannot be arbitrarily small (minimum feature criteria for dihedral angles) there exists a certain translation such that the distance between v_1' and v_2' will be $> \delta$. Set the minimum dihedral angle α_{min} between any two facets to be such that $\frac{\delta}{\sin\alpha_{min}} \leq \frac{M}{2}$. By Lemma 4.1, the ordering of v_1, v_2 on L which is consistent with the classification of v can be determined exactly. The ambiguity in the ordering of *old vertices* and *new vertices* on the edges which are not incident on a common vertex does not arise if we assume minimum feature separation of at least δ for elements of the input polyhedron S.

Generating new edges: Let L be the line of intersection of a facet f with the *notch plane*. Let $(v_1, v_2, ..., v_k)$ be the sorted sequence of vertices on L, corresponding to the points of intersection between the facet and the *notch plane*. One needs to decide consistently whether there should be an edge between two consecutive vertices v_i and v_{i+1} of this sorted sequence. This is done by scanning these sorted vertices from one end to the other and deciding whether we are "inside" or "outside" the facet. It is easy to see that if v_i is a *new vertex* then there would be an edge between v_i and v_{i+1} if there were no edge between v_{i-1} and v_i and vice versa. But if v_i is an *old vertex* there can be edge between v_i and v_{i+1} disregard of the presence of an edge between v_{i-1}, v_i.
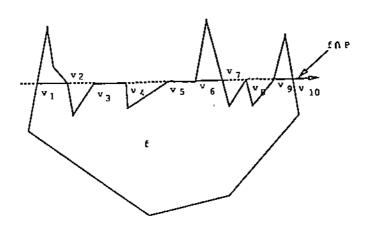

Figure 11: Generating *new* and *old* edges.

Toggling between "inside" and "outside" of the facet is carried out properly, even with degeneracies, using a *multiplicity code* at each intersection vertex. Scan the sorted sequence of intersection vertices from one end to the other and maintain a counter which is incremented by the *multiplicity code* at each vertex. Toggle between "inside" and "outside" of the facet as the counter toggles between "odd" and "even" count. For a *new vertex* put a *multiplicity code* of 1. For an *old vertex*, put a *multiplicity code* of 1 if two incident edges on the vertex on that facet lie in different half-spaces of P_g and put a *multiplicity code* of 2 if they lie in the same half-space. If there is an *old edge* between two vertices v_i and v_{i+1}, put *multiplicity codes* on them as follows. If other two incident edges on v_i, v_{i+1} on the facet f lie in the same half-space of the *notch plane*, put a *multiplicity code* of 1 on both the vertices v_i and v_{i+1}. Otherwise, put *multiplicity codes* of 1 and 2 on v_i and v_{i+1} in any order. In Figure 11, there is an *old edge* between v_3, v_4. The status ("outside") with which one enters the vertex v_3 is same as that one with which one leaves the vertex v_4. This is enforced by putting a *multiplicity code* of 1 on the two vertices which increment the counter by an "even" amount and prevent it from toggling. There is another *old edge* between v_5 and

v_6. The status ("outside") with which one enters the vertex v_5 is different from the one with which one leaves the vertex v_6. This is enforced by putting *multiplicity codes* of 1 and 2 on the two vertices in any order which increment the counter by an "odd" amount and make it toggle. Initially, the counter is set to 0. Create a *new edge* from vertex v_i to v_{i+1} if the count is "odd" after leaving the vertex v_i. In case, there is an *old edge* between v_i and v_{i+1}, skip creating any *new edge* between them. An *old edge* may be transferred to GP_g^ℓ or GP_g^r or to both. Transferring of *old edges* is described below.

Transfer of old edges: The *old edge* e_o should be transferred to GP_g^ℓ (GP_g^r respectively.) if any facet (or a part of it) adjacent to e_o which has not been decided to be on the *notch plane*, gets transferred to GP_g^ℓ (GP_g^r respectively.). For example, the edge g in Figure 3 should be transferred to GP_g^ℓ but not to GP_g^r. For each *old edge* e_o decided to be on the plane P_g, check all of its *oriented edges* on different facets which have not been decided to be on the *notch plane*. Suppose f_o is such a facet. Classify any vertex v_o of f_o w.r.t the oriented edge e_o on f_o. If it is on the same side of e_o in which f_o lies then e_o should be transferred to GP_ℓ (GP_r respectively.) if v_o has been classified to lie in P^+ (P^- respectively.). It is trivial to decide the side of e_o in which f_o lies from the *oriented edge* of e_o on f_o.

Consistent vertex-plane, edge-plane and facet-plane classification takes overall $O(p)$ time where p is the total number edges of the polyhedron S. The above bound follows from the fact that each edge of S is visited only $O(1)$ time to determine the intersection points of S with the *notch plane* P_g. The sorting of intersection points on the facets adds $O(r \log r + q')$ time where q' is the total number of facets decided to be intersected by the *notch plane*. The above bound follows from the fact that any line segment intersects a facet having r_i *reflex vertices* in no more than $(2r_i + 2)$ points (Lemma 2.4). Once the construction of the maps GP_g^ℓ and GP_g^r is done, it is trivial to recognize the boundary B_g containing the *notch* g. The methods as described in section 3 can be used to determine the *interesting boundaries*. Note that, if B_g is an inner boundary, the *interesting boundaries* consist of all the *ancestors* of B_g. If all the polygons in the *cross sectional maps* are *fleshy*, ancestors of B_g can be determined exactly using Lemma 3.4 of [3] at a cost of $O(u')$ where u' is the number of vertices on the *cross sectional maps*. As discussed earlier, there are $O(t)$ polygons and *monotone chains* in the *interesting boundaries* where t is the number of *notches* intersected by P_g. Let u be the number of vertices on the *interesting boundaries*. According to Lemma 4.2, the children and parent of B_g can be determined exactly in $O(t^2 + u(t + \log u))$ time if the polygons corresponding to the *interesting boundaries* are *fleshy*. Set up a safe minimum feature separation between polyhedral features so that the polygons generated in the *cross sectional maps* are always *fleshy*. Detection of children and parent of the polygon containing the *notch* g in effect, determines the inner and outer boundaries of $Q_g^\ell(Q_g^r)$. Obviously $q' = O(p)$ and $u = O(u') = O(p)$. Combining the complexities of computing the edges of GP_g^ℓ (GP_g^r respectively.) and detecting the inner and outer boundaries of Q_g^ℓ (Q_g^r respectively.), we conclude that Q_g^ℓ (Q_g^r respectively.) can be computed in $O(p + t^2 + u'(t + \log u'))$ time.

Step II: S is separated corresponding to the cut $Q_g^l(Q_g^r)$ by splitting the facets which are intersected by the *cut* Q_g. Let f_i be such a facet which is to be split at $a_1, a_2, ..., a_k$. For each such point of intersection which may correspond to a *new vertex* or an *old vertex*, do the following.

New Vertex: Let $e_s = (v_1, v_2)$ be the edge on which *new vertex* v_n lies. Generate edges between v_1, v_n and between v_2, v_n. Since the half spaces in which v_1 and v_2 lie are known, one can decide the half space in which each such *new edges* lies.

Old Vertex: For each *old vertex* v_o lying on the plane P_g, transfer the edges connected to v_o to the half

23

space in which their other vertex has been decided to lie in. Here, transferring means connecting those edges to the copy of the vertex v_o on the corresponding cut. The edges connected to v_o which have been decided to be on the plane P_g are transferred by procedure as described before. Finally, create two facets corresponding to the cuts Q_g^ℓ and Q_g^r. Splitting each facet which are decided to be intersected by the cut Q_g^ℓ (Q_g^r) effectively either splits S into separate pieces or splices it about the cuts creating two facets corresponding to the cuts at the same geometric location. A depth first search starting from one vertex in each of P_g^+ and P_g^- resolves this ambiguity and also collects all the features pertinent to each piece. Certainly, this separation step does not take more than $O(p)$ time where p is the number of edges of S. Combining the time and space complexities of *Step I* and *Step II* we have the following Lemma.

Lemma 4.3 Using heuristics to avoid conflicting decisions, a manifold polyhedron S with arbitrary genus, shells and certain minimum feature separations can be partitioned under finite precision arithmetic with a *notch plane* in $O(p + t^2 + u'(t + logu'))$ time and $O(p)$ space, where p is the number of edges in S, u' is the number of vertices on the *cross sectional maps* and t is the number of *notches* intersected by the *notch plane*.

The following combinatorial Lemma is used to derive the time complexity in Theorem 4.1.
Lemma 4.4. Let $S_1, S_2, ..., S_k$ be the polyhedra in the current decomposition, where each S_i contains a subnotch g_i of a notch g of a manifold polyhedron S with n edges and r notches, and let u_i' be the total number of vertices on the *cross sectional map* in S_i. Then we have $u' = \sum_{i=1}^{k} u_i{}' = O(n + r^2)$, where u' is the total number of vertices on the *cross sectional maps* in $S_1, S_2, ..., S_k$.

Proof: Consider the *cross sectional map* GP_g^ℓ (GP_g^r). The lines of intersection between P_g and other *notch planes*, called the *notch lines* divide this map into smaller facets which are present on the *cross sectional maps* in $S_1, S_2, ..., S_k$ i.e. on $\cup_{i=1}^{k} GP_{g_i}^\ell$ ($\cup_{i=1}^{k} GP_{g_i}^r$). The vertices on $\cup_{i=1}^{k} GP_{g_i}^\ell$ ($\cup_{i=1}^{k} GP_{g_i}^r$) can be partitioned into three sets, viz., T_1, T_2 and T_3. The set T_1 consists of vertices which are created by intersections between *notch lines*. The set T_2 consists of vertices on GP_g^ℓ (GP_g^r) and the set T_3 consists of vertices which are created by intersections between edges of GP_g^ℓ (GP_g^r) and *notch lines*. Since there are at most $O(r)$ *notch lines*, $|T_1| \leq r^2$. Certainly, $|T_2| \leq n$. By Lemma 2.4, each *notch line* can intersect GP_g^ℓ (GP_g^r) in at most $(2r+2)$ points since GP_g^ℓ (GP_g^r) can have at most $O(r)$ *reflex vertices*. This gives $|T_3| \leq 2r + 2$. Thus,

$$
\begin{aligned}
u' = \sum_{i=1}^{k} u_i' &= |T_1| + |T_2| + |T_3| \\
&\leq r^2 + n + 2r + 2 \\
&= O(n + r^2). \clubsuit
\end{aligned}
$$

Theorem 4.1 Using heuristics to avoid conflicting decisions, a polyhedron S with arbitrary number of holes and shells and certain minimum feature separations can be decomposed under finite precision arithmetic into $O(r^2)$ convex pieces in $O(nr^2 + nrlogn + r^3logn + r^4)$ time and in $O(nr)$ space, where r is the number of *notches*, n is the number of edges in S.

Proof: Let S be a manifold polyhedron. At a generic instance of the algorithm, let $S_1, S_2, ..., S_k$ be the

24

k distinct (non-convex) polyhedra in the current decomposition, which contain the *subnotches* of a *notch* g which is to be removed. Let p_i be the number of edges in S_i, u_i' be the number of vertices on the *cross sectional maps* in S_i and t_i be the number of *notches* intersected by the *notch plane* in S_i. Let $p = \sum_{i=1}^{k} p_i$, $u' = \sum_{i=1}^{k} u_i'$ and $t = \sum_{i=1}^{k} t_i$. Certainly, $k = O(r)$ and $t = O(r)$. Using Lemma 4.3, we can say that the time $\Im$ to remove the *notch* g is given by

$$\Im = O(\sum_{i=1}^{k}(p_i + t_i^2 + u_i'(t_i + log u_i')))$$
$$= O(p + r^3 + u'r + u'log u')$$

By Lemma 4.4, $u' = O(n + r^2)$. This gives,

$$\Im = O(p + r^2 + (n + r^2)r + (n + r^2)log n)$$
$$= O(nr + nlog n + r^2 log n + r^3)$$

To carry out removal of r *notches* we need $O(nr^2 + nrlog n + r^3 log n + r^4)$ time. Obviously, the space complexity is $O(p) = O(nr)$. If S is a non-manifold polyhedron, remove all *special notches* from S to produce manifold polyhedra and decompose each such polyhedron into convex pieces as discussed in the previous section. The complexity remains same for this case. ♣

5 Conclusion

We have implemented our polyhedral decomposition algorithm under floating point arithmetic in Common Lisp on a Symbolics 3650. The numerical computations are all in C, callable from Lisp. We used $\delta = 2^{-17}$ in the 32 bit machine with precision 2^{-25}. Simple examples are shown in Figure 12. The experimental results have been very satisfying. Test polyhedra were generated by SHILP solid model creation software.

Our next goal is to develop a robust and stable algorithm for polyhedral decomposition problem. To find a robust and stable algorithm for this problem seems to be quite hard. It may be worthwhile to consider the concept of *pseudo facets*, the counterpart of *pseudo lines* in three dimensions to solve this problem.
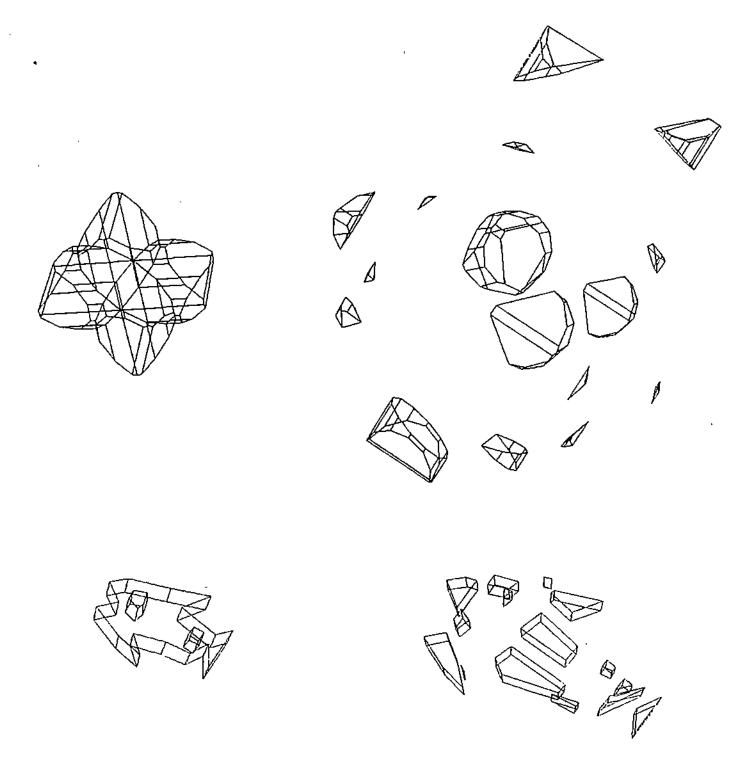
Figure 12: Examples.

References

[1] Bajaj, C., (1989) "Geometric Modeling with Algebraic Surfaces", *The Mathematics of Surfaces III*, edited by D. Handscomb, Oxford University Press, 3 - 48.

[2] Anupam, V., Bajaj, C., Dey, T., Fields, M., Ihm, I., Klinkner, S., (1989), "The SHILP solid model creation, editing and display toolkit", Manuscript.

[3] Bajaj, C., and Dey, T., (1989) "Polygon Nesting and Robustness" *Proc. Intl. Workshop on Discrete Algorithms and Complexity*, Fukuoka, Japan, 33-40. To appear in *Information Processing Letters*.

[4] Chazelle, B., (1980), "Computational Geometry and Convexity", *Ph.D. Thesis*, CMU-CS-80-150, Computer Science, Carnegie-Mellon University.

[5] Chazelle, B., (1984), "Convex Partitions of Polyhedra: A Lower Bound and Worst-case Optimal Algorithm", *SIAM J. on Computing*, Vol. 13, No. 3, pp. 488–507.

[6] Chazelle, B., and Palios, L., (1989), "Triangulating a Non-convex Polytope" *Proc. of the 5th ACM Symposium on Computational Geometry*, Saarbrucken, West Germany, 393-400.

[7] Dobkin, D., and Silver, D., (1988), "Recipes for Geometry and Numerical Analysis", *Proc. of the Fourth ACM Symposium on Computational Geometry*, Urbana, Illinois, 93 - 105.

[8] Edelsbrunner, H., (1987), "Algorithms in Combinatorial Geometry", Springer Verlag.

[9] Edelsbrunner, H., and Mucke, P., (1988), "Simulation of Simplicity: A Technique to Cope with Degenerate Cases in Geometric Algorithms" *Proc. of the Fourth ACM Symposium on Computational Geometry*, Urbana, Illinois, 118-133.

[10] Fortune, S., (1989) "Stable Maintenance of Point-set Triangulations in Two Dimensions", *Proc. 30th IEEE Symposium on the Foundations of Computer Science*, 494 - 499.

[11] Guibas, L., Salesin, D., and Stolfi, J., (1989) "Building Robust Algorithms from Imprecise Computations", *Proc. 1989 ACM Symposium on Computational Geometry*, Saarbuchen, West Germany, 208-217.

[12] Hoffmann, C., Hopcroft, J., and Karasick, M., (1987), "Robust Set Operations on Polyhedral Solids", Dept. of Computer Science, Cornell University, Technical Report 87-875.

[13] Hoffmann, K., Mehlhorn, K., Rosenstiehl, P., and Tarjan, R., (1986), "Sorting Jordan Sequences in Linear Time using Level Linked Search Trees", *Information and Control*, 68, 170 - 184.

[14] Hopcroft, J., Kahn, P., "A Paradigm for Robust Geometric Algorithms", Computer Science Tech. Report, Cornell University, TR 89-1044.

[15] Karasick, M., (1988) "On the Representation and Manipulation of Rigid Solids", Ph.D. Thesis. McGill University.

[16] Lingas, A., (1982), "The Power of Non-Rectilinear Holes", *Proc. 9th Intl. Colloquium on Automata, Languages and Programming*, Lecture Notes in Computer Science, Springer Verlag, 369 - 383.

[17] Milenkovic, V., (1988), "Verifiable Implementations of Geometric Algorithms Using Finite Precision Arithmetic", Ph.D. Thesis, CMU Tech. Report CS-88-168, Carnegie Mellon Univ., Pittsburgh.

[18] O'Rourke, J., and Supowit, K., (1983), "Some NP-hard Polygon Decomposition Problems". *IEEE Trans. Inform. Theory*, 29, 181 - 190.

[19] Requicha, A.A.G., (1977), "Mathematical Models of Rigid Solid Objects", Tech. Memo 28, Production Automation Project, University of Rochester, Rochester, NY.

[20] Rupert, J., and Seidel, R., (1989) "On the Difficulty of Tetrahedralizing Three Dimensional Nonconvex Polyhedra", *Proc. of the Fifth ACM Symposium on Computational Geometry*, Saarbrucken, West Germany, 380 - 392.

[21] Segal, M., and Sequin, C.. (1985), "Consistent Calculations for Solids Modeling", *Proc. of the First ACM Symposium on Computational Geometry*, 29 - 38.

[22] Sugihara, K., (1988), "A Simple Method of Avoiding Numerical errors and Degeneracy in Voronoi diagram Constructions", Research Memorandum RMI 88-14, Department of Mathematical Engineering and Instrumentation Physics, Tokyo University.

[23] Sugihara, K., and Iri, M.. (1989), "A Solid Modeling System Free from Topological Consistency", Research Memorandum RMI 89-3, Department of Mathematical Engineering and Instrumentation Physics, Tokyo University.

[24] Yap, C.. (1988) "A Geometric Consistency Theorem for a Symbolic Perturbation Theorem" *Proc. of the Fourth ACM Symposium on Computational Geometry*, Urbana, Illinois, 134-142.